

```
In [1]: import numpy as np
import pandas as pd
import os
os.chdir("C:\\Users\\kumar\\OneDrive\\Desktop\\Machine Learning")
dataset = pd.read_csv('Churn_Modelling.csv')
dataset.head()
```

```
Out[1]:
```

|   | RowNumber | CustomerId | Surname  | CreditScore | Geography | Gender | Age | Tenure | Balance   |
|---|-----------|------------|----------|-------------|-----------|--------|-----|--------|-----------|
| 0 | 1         | 15634602   | Hargrave | 619         | France    | Female | 42  | 2      | 0.00      |
| 1 | 2         | 15647311   | Hill     | 608         | Spain     | Female | 41  | 1      | 83807.86  |
| 2 | 3         | 15619304   | Onio     | 502         | France    | Female | 42  | 8      | 159660.80 |
| 3 | 4         | 15701354   | Boni     | 699         | France    | Female | 39  | 1      | 0.00      |
| 4 | 5         | 15737888   | Mitchell | 850         | Spain     | Female | 43  | 2      | 125510.82 |

```
In [2]: dataset.columns
```

```
Out[2]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
            'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
            'IsActiveMember', 'EstimatedSalary', 'Exited'],
            dtype='object')
```

```
In [3]: dataset['Geography'].value_counts()
```

```
Out[3]: France      5014
Germany    2509
Spain      2477
Name: Geography, dtype: int64
```


```
In [4]: print(dataset.groupby('Geography')['CreditScore'].mean())
```

```
Geography
France      649.668329
Germany     651.453567
Spain       651.333872
Name: CreditScore, dtype: float64
```

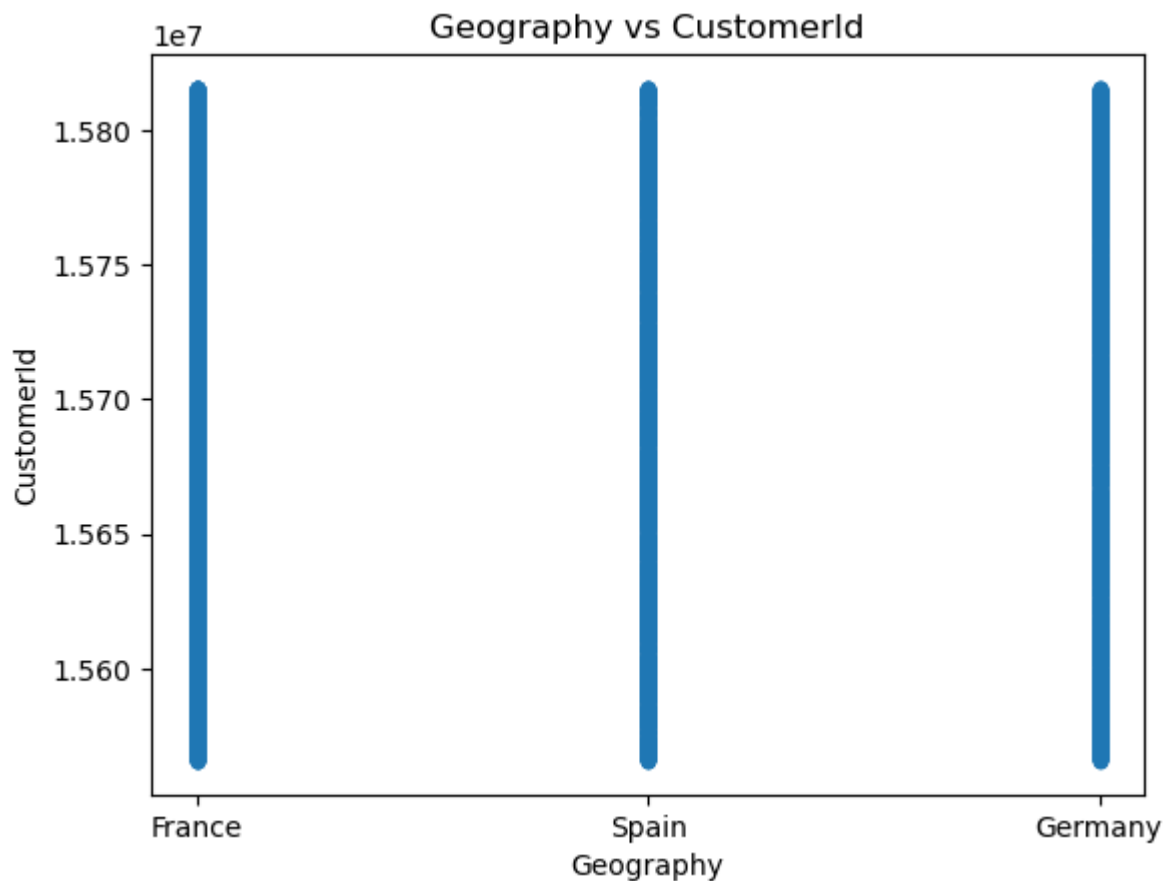
```
In [5]: print(dataset.groupby('CustomerId')['Geography'].value_counts())
```

```
CustomerId  Geography
15565701    Spain      1
15565706    Spain      1
15565714    France     1
15565779    Germany    1
15565796    Germany    1
..
15815628    France     1
15815645    France     1
15815656    Germany    1
15815660    France     1
15815690    Spain      1
Name: Geography, Length: 10000, dtype: int64
```

```
In [6]: import matplotlib.pyplot as plt
import seaborn as sns
plt.hist(dataset['EstimatedSalary'], bins = 100)
plt.show()
```



```
In [7]: dataset.plot.scatter(x="Geography",y="CustomerId")
plt.title("Geography vs CustomerId")
plt.xlabel("Geography")
plt.ylabel("CustomerId")
plt.show()
```

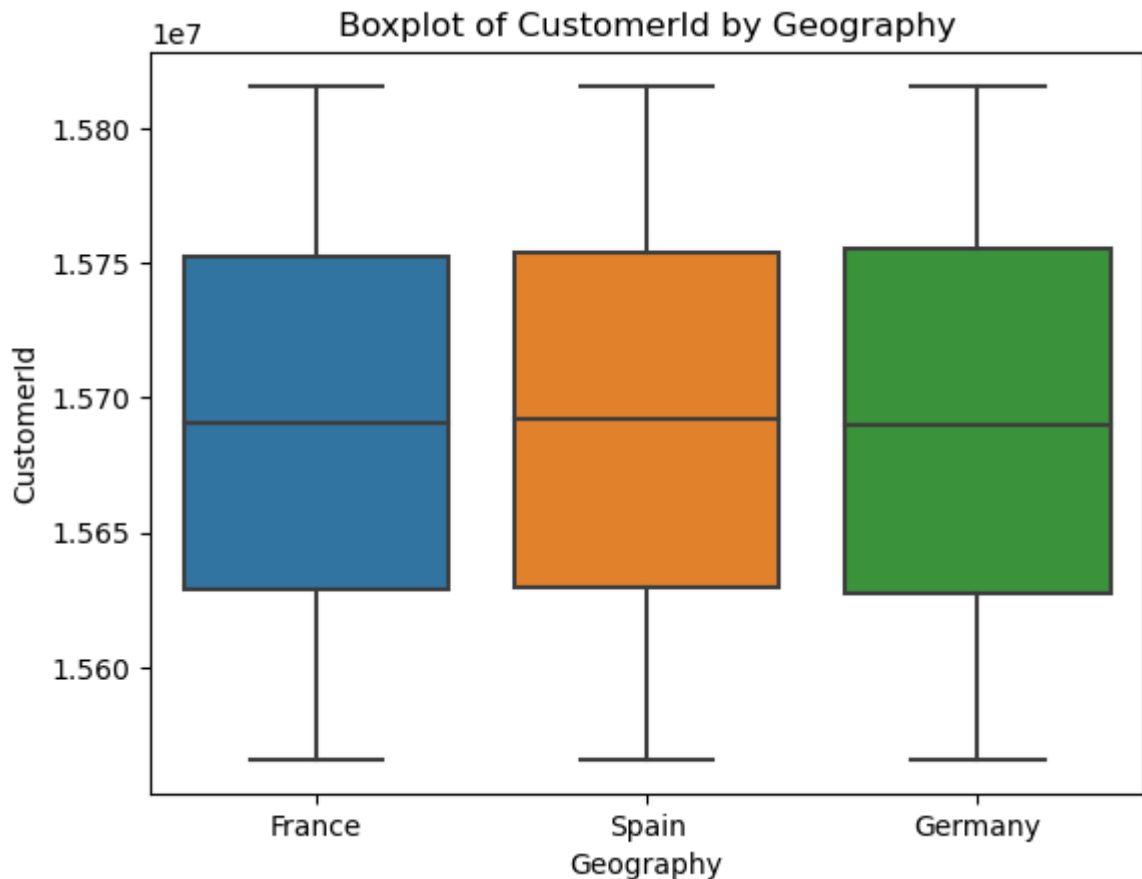


```
In [8]: import seaborn as sns
import matplotlib.pyplot as plt
```

```
sns.boxplot(x='Geography', y='CustomerId', data=dataset)

plt.title('Boxplot of CustomerId by Geography')
plt.xlabel('Geography')
plt.ylabel('CustomerId')

plt.show()
```



```
In [9]: print(dataset.shape)

(10000, 14)
```

```
In [10]: X = dataset.iloc[:, 0:13].values # Features (columns 0 to 12)
y = dataset.iloc[:, 13].values # Target variable (column 13)
from sklearn.preprocessing import LabelEncoder
labelencoder_X_1 = LabelEncoder()
X[:, 3] = labelencoder_X_1.fit_transform(X[:, 3])
labelencoder_X_2 = LabelEncoder()
X[:, 4] = labelencoder_X_2.fit_transform(X[:, 4])
labelencoder_y = LabelEncoder()
y = labelencoder_y.fit_transform(y)
```

```
In [11]: categorical_column = dataset.iloc[:, 0]
X_State = pd.get_dummies(categorical_column, drop_first=True)
X_State = X_State.astype(int)
dataset = dataset.drop(columns=[categorical_column.name])
X = pd.concat([X_State, dataset], axis=1)
```

```
In [12]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=42)
```

```
In [13]: from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
```

```
In [14]: dataset = dataset.drop(['CustomerId', 'Surname'], axis=1)
categorical_columns = ['Gender', 'Geography']
ct = ColumnTransformer(
    transformers=[('encoder', OneHotEncoder(), categorical_columns)],
    remainder='passthrough'
)
X = ct.fit_transform(dataset.drop('Exited', axis=1))
y = dataset['Exited']
```

```
In [15]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [16]: sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
In [17]: clf = RandomForestClassifier()
clf.fit(X_train, y_train)
```

```
Out[17]: RandomForestClassifier()
```

```
In [18]: y_pred = clf.predict(X_test)
```

```
In [19]: from sklearn.metrics import accuracy_score
```

```
In [20]: accuracy_score(y_test, y_pred)
```

```
Out[20]: 0.8635
```

```
In [21]: from sklearn.metrics import confusion_matrix
print(confusion_matrix(y_test, y_pred))
```

```
[[1544  63]
 [ 210 183]]
```

## Result Observed

True Positives(TP) = 183

True Negatives(TN) = 1544

False Positives(FP) = 63

False Negatives(FN) = 210

Precision =  $TP / (TP + FP) = 183 / (183 + 63) = 183 / 246 = 0.744$

$$\text{Recall} = \text{TP}/(\text{TP}+\text{FN}) = 183/(183+210)=183/393=0.4656$$

$$\text{Accuracy} = (\text{TP}+\text{TN})/(\text{TP}+\text{TN}+\text{FP}+\text{FN}) = (183+1544)/(183+1544+63+210)=1727/2000=0.8635$$

In [ ]: