

```
In [1]: import numpy as np
import pandas as pd
import warnings
import os
os.chdir("C:\\Users\\kumar\\OneDrive\\Desktop\\Machine Learning")
def warn(*args, **kwargs):
    pass
warnings.warn = warn
```

```
In [2]: df = pd.read_csv("spam.csv", encoding='ISO-8859-1')
df.head()
```

```
Out[2]:
```

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	NaN	NaN

```
In [3]: len(df)
```

```
Out[3]: 5572
```

```
In [4]: df.isnull().sum()
```

```
Out[4]: v1          0
v2          0
Unnamed: 2    5522
Unnamed: 3    5560
Unnamed: 4    5566
dtype: int64
```

```
In [5]: df = df[['v1', 'v2']]
df.head()
```

```
Out[5]:
```

	v1	v2
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

```
In [6]: blanks = []
for i, v1, v2 in df.itertuples():
    if type(v2) == str and v2.isspace():
        blanks.append(i)
if len(blanks) > 0:
    print(f'There are len(blanks) empty space strings in the dataset')
    df.drop(blanks, inplace=True)
```

```
In [7]: df.nunique()
```

```
Out[7]: v1      2
        v2    5169
        dtype: int64
```

```
In [8]: df['v1'].value_counts()
```

```
Out[8]: ham      4825
        spam     747
        Name: v1, dtype: int64
```

```
In [9]: from scipy import stats

df['v2_length'] = df['v2'].apply(len)

hams = df[df['v1'] == 'ham']['v2_length']
spams = df[df['v1'] == 'spam']['v2_length']

t_stat, p_value = stats.ttest_ind(hams, spams)

alpha = 0.05

if p_value < alpha:
    print('The difference in sms lengths are significance')
else:
    print('The difference in sms lengths are not significance')

print(f'T-Stat: {t_stat}')
print(f'P-Value: {p_value}')
```

```
The difference in sms lengths are significance
T-Stat: -31.350650338992136
P-Value: 7.702078585492358e-199
```

```
In [10]: X = df['v2']
        y = df['v1']
        display(X, y)
```

```
0      Go until jurong point, crazy.. Available only ...
1              Ok lar... Joking wif u oni...
2      Free entry in 2 a wkly comp to win FA Cup fina...
3      U dun say so early hor... U c already then say...
4      Nah I don't think he goes to usf, he lives aro...

...
5567    This is the 2nd time we have tried 2 contact u...
5568              Will i_b going to esplanade fr home?
5569    Pity, * was in mood for that. So...any other s...
5570    The guy did some bitching but I acted like i'd...
5571              Rofl. Its true to its name
Name: v2, Length: 5572, dtype: object
0      ham
1      ham
2      spam
3      ham
4      ham

...
5567    spam
5568      ham
5569      ham
5570      ham
5571      ham
Name: v1, Length: 5572, dtype: object
```

```
In [11]: from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()
```

```

y = le.fit_transform(y)

display(y)

array([0, 0, 1, ..., 0, 0, 0])

```

```

In [12]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_sta

```

```

In [13]: from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.svm import LinearSVC

```

```

In [14]: from sklearn.model_selection import GridSearchCV

pipeline = Pipeline([('tfidf', TfidfVectorizer()), ('classifier', LinearSVC())])

grid_params_svm = {
    'tfidf__ngram_range': [(1,1), (1,2)],
    'tfidf__stop_words': [None, 'english'],
    'classifier__C': [0.01, 0.1, 0.5, 1.0, 1.5, 2.0, 2.5, 5.0, 10.0],
}

model_svm = GridSearchCV(pipeline, grid_params_svm, cv=5, scoring='accuracy')

model_svm.fit(X_train, y_train)

```

```

Out[14]: GridSearchCV(cv=5,
    estimator=Pipeline(steps=[('tfidf', TfidfVectorizer()),
    ('classifier', LinearSVC())]),
    param_grid={'classifier__C': [0.01, 0.1, 0.5, 1.0, 1.5, 2.0, 2.5,
    5.0, 10.0],
    'tfidf__ngram_range': [(1, 1), (1, 2)],
    'tfidf__stop_words': [None, 'english']},
    scoring='accuracy')

```

```

In [15]: from sklearn.naive_bayes import MultinomialNB

pipeline = Pipeline([('tfidf', TfidfVectorizer()), ('classifier', MultinomialNB())])

grid_params_nb = {
    'tfidf__ngram_range': [(1,1), (1,2)],
    'tfidf__stop_words': [None, 'english'],
    'classifier__alpha': [0.01, 0.1, 0.5, 1.0, 1.5, 2.0, 2.5, 5.0, 10.0]
}

model_nb = GridSearchCV(pipeline, grid_params_nb, cv=5, scoring='accuracy')

model_nb.fit(X_train, y_train)

```

```

Out[15]: GridSearchCV(cv=5,
    estimator=Pipeline(steps=[('tfidf', TfidfVectorizer()),
    ('classifier', MultinomialNB())]),
    param_grid={'classifier__alpha': [0.01, 0.1, 0.5, 1.0, 1.5, 2.0,
    2.5, 5.0, 10.0],
    'tfidf__ngram_range': [(1, 1), (1, 2)],
    'tfidf__stop_words': [None, 'english']},
    scoring='accuracy')

```

```

In [16]: def show_metrics(y_true, y_pred, grid_search=None):
    from sklearn.metrics import (classification_report,
    confusion_matrix,
    ConfusionMatrixDisplay)

```

```

print('-' * 20)
print(classification_report(y_true, y_pred))
print(confusion_matrix(y_true, y_pred))

if grid_search:
    print('-' * 20)
    print(grid_search.best_params_)

```

```

In [17]: best_svm = model_svm.best_estimator_
y_pred_svm = best_svm.predict(X_test)

show_metrics(y_test, y_pred_svm, model_svm)

```

```

-----
              precision    recall  f1-score   support

     0       0.99         1.00         0.99         1464
     1       0.97         0.93         0.95          208

 accuracy          0.99         0.99         0.99         1672
 macro avg         0.98         0.96         0.97         1672
 weighted avg         0.99         0.99         0.99         1672

```

```

[[1458    6]
 [  14 194]]

```

```

-----
{'classifier__C': 10.0, 'tfidf__ngram_range': (1, 2), 'tfidf__stop_words': None}

```

```

In [18]: best_nb = model_nb.best_estimator_
y_pred_nb = best_nb.predict(X_test)

show_metrics(y_test, y_pred_nb, model_nb)

```

```

-----
              precision    recall  f1-score   support

     0       0.99         0.99         0.99         1464
     1       0.96         0.92         0.94          208

 accuracy          0.99         0.99         0.99         1672
 macro avg         0.97         0.96         0.97         1672
 weighted avg         0.99         0.99         0.99         1672

```

```

[[1456    8]
 [  16 192]]

```

```

-----
{'classifier__alpha': 0.01, 'tfidf__ngram_range': (1, 2), 'tfidf__stop_words': None}

```

```

In [19]: def save_model(model, prefix=''):
import joblib
from datetime import datetime

# Get the current date and time as a string to define the file name
current_datetime = datetime.now().strftime("%Y-%m-%d_%H-%M-%S")
filename = f"{prefix}model_{current_datetime}.joblib"

joblib.dump(model, filename)

print(f"Model saved to {filename}")

```

```

In [20]: save_model(model_svm, prefix='svm_')
save_model(model_nb, prefix='nb_')

```

